

hw5q3

November 18, 2024

1 set-up

```
[1]: !pip3 install scikit-learn
# import statements
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
from sklearn.datasets import fetch_openml
!pip install --upgrade certifi
!brew install openssl
!brew link openssl --force
import certifi
import ssl
import os

os.environ['SSL_CERT_FILE'] = certifi.where()
import requests

url = 'https://example.com'
response = requests.get(url, verify=False)
!/Applications/Python\ 3.12.4/Install\ Certificates.command
import ssl
ssl._create_default_https_context = ssl._create_unverified_context
# this cell will take a minute to run depending on your internet connection
X, y = fetch_openml('mnist_784', version=1, return_X_y=True) # getting data
    ↪ from online

print('X shape:', X.shape, 'y shape:', y.shape)
# this cell processes some of the data

# if this returns an error of the form "KeyError: 0", then try running the
    ↪ following first:
X = X.values # this converts X from a pandas dataframe to a numpy array

digits = {j:[] for j in range(10)}
for j in range(len(y)): # takes data assigns it into a dictionary
    digits[int(y[j])].append(X[j].reshape(28,28))
```

```

digits = {j:np.stack(digits[j]) for j in range(10)} # stack everything to be
↳one numpy array
for j in range(10):
    print('Shape of data with label', j, ':', digits[j].shape )
    # this cell would stack 100 examples from each class together
# this cell also ensures that each pixel is a float between 0 and 1 instead of
↳an int between 0 and 255
data = []
for i in range(10):
    flattened_images = digits[i][:100].reshape(100,-1)
    data.append(flattened_images)

data = np.vstack(data)
data = data.astype('float32') / 255.0

```

Requirement already satisfied: scikit-learn in
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages
(1.5.2)

Requirement already satisfied: numpy>=1.19.5 in
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages
(from scikit-learn) (1.26.4)

Requirement already satisfied: scipy>=1.6.0 in
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages
(from scikit-learn) (1.14.1)

Requirement already satisfied: joblib>=1.2.0 in
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages
(from scikit-learn) (1.4.2)

Requirement already satisfied: threadpoolctl>=3.1.0 in
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages
(from scikit-learn) (3.5.0)

Requirement already satisfied: certifi in
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages
(2024.8.30)

=> Downloading https://formulae.brew.sh/api/formula.jws.json

100.0%

=> Downloading https://formulae.brew.sh/api/cask.jws.json

100.0%

Warning: openssl@3 3.4.0 is already installed and up-to-date.

To reinstall 3.4.0, run:

```
brew reinstall openssl@3
```

Warning: Already linked: /opt/homebrew/Cellar/openssl@3/3.4.0

To relink, run:

```
brew unlink openssl@3 && brew link openssl@3
```

zsh:1: no such file or directory: /Applications/Python 3.12.4/Install

Certificates.command

/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-
packages/urllib3/connectionpool.py:1099: InsecureRequestWarning: Unverified

HTTPS request is being made to host 'example.com'. Adding certificate verification is strongly advised. See:
<https://urllib3.readthedocs.io/en/latest/advanced-usage.html#tls-warnings>
warnings.warn(

```
X shape: (70000, 784) y shape: (70000,)
Shape of data with label 0 : (6903, 28, 28)
Shape of data with label 1 : (7877, 28, 28)
Shape of data with label 2 : (6990, 28, 28)
Shape of data with label 3 : (7141, 28, 28)
Shape of data with label 4 : (6824, 28, 28)
Shape of data with label 5 : (6313, 28, 28)
Shape of data with label 6 : (6876, 28, 28)
Shape of data with label 7 : (7293, 28, 28)
Shape of data with label 8 : (6825, 28, 28)
Shape of data with label 9 : (6958, 28, 28)
```

2 Extract 4s and 9s

```
[2]: # create dataset here (essentially just create a numpy array of 1's and -1's
      ↪ for the labels)

#zeros = digits[0][:500]
#ones = digits[1][:500]
nbr_examples = 500
numbers = [4,9]
map = {((-1)**(i+1)): [digit.flatten() for digit in digits[i][:nbr_examples]]
      ↪ for i in numbers}
for key in [-1, 1]:
    x = map[key]
    for i in range(len(x)):
        x_i = (x[i] - x[i].min()) / (x[i].max() - x[i].min())
        map[key][i] = x_i
```

3 Function and gradients....

```
[3]: def function(weights:[float], map:{int:np.array}) -> float:
      sum = 0
      for i in [-1,1]:
          for x_i in map[i]:
              exponent = float(np.dot(weights, x_i)) * i * (-1)
              inner_expression = 1 + math.exp(exponent)
              sum += math.log(inner_expression)
      return sum / (len(map[-1]) + len(map[1]))
```

```

def gradient(weights:[float], map:{int:np.array}) -> np.array:
    gradient_array = np.zeros(len(weights))
    N = len(weights)
    for w_index in range(N):
        for i in [-1,1]:
            for x_i in map[i]:
                exponent = float(np.dot(weights, x_i)) * i * (-1)
                numerator = (-i) * x_i[w_index] * math.exp(exponent)
                denominator = 1 + math.exp(exponent)
                gradient_array[w_index] += numerator / denominator
            gradient_array[w_index] /= N
    return gradient_array

def optimized_gradient(weights:[float], map:{int:np.array}) -> np.array:
    gradient_array = np.zeros(len(weights))
    N = len(weights)
    for i in [-1,1]:
        for x_i in map[i]:
            exponent = float(np.dot(weights, x_i)) * i * (-1)
            numerator = (-i) * math.exp(exponent)
            denominator = 1 + math.exp(exponent)
            for w_index in range(N):
                gradient_array[w_index] += numerator * x_i[w_index] /
    ↪denominator
    return np.array([grad / N for grad in gradient_array])

def printf(weights:[float], map:{int:np.array}, i:int):
    if (i % 1000) == 0:
        print(f"iteration : {i}, F(w) = {function(weights, map)}")
    #else:
    #    pass
    #print(f"iteration : {i}, F(w) = {function(weights, map)}")

def get_P(weights:[float], map:{int:np.array}):
    fw = optimized_gradient(weights, map)
    return np.sign(fw) * np.sum(np.abs(fw))

def get_P1(weights: np.array, map: {int: np.array}, count_vector):
    fw = optimized_gradient(weights, map)
    j_star = np.argmax(np.abs(fw))
    if count_vector is not None:
        count_vector[j_star] += 1
    p_t = np.zeros_like(fw)
    p_t[j_star] = np.sign(fw[j_star]) * np.max(np.abs(fw))

    return p_t

```

4 Backtracking line search

```
[38]: # 4 a
import math
iteration_nbr = 10**4

weights = np.array([0 for _ in range(784)], dtype=np.float64)

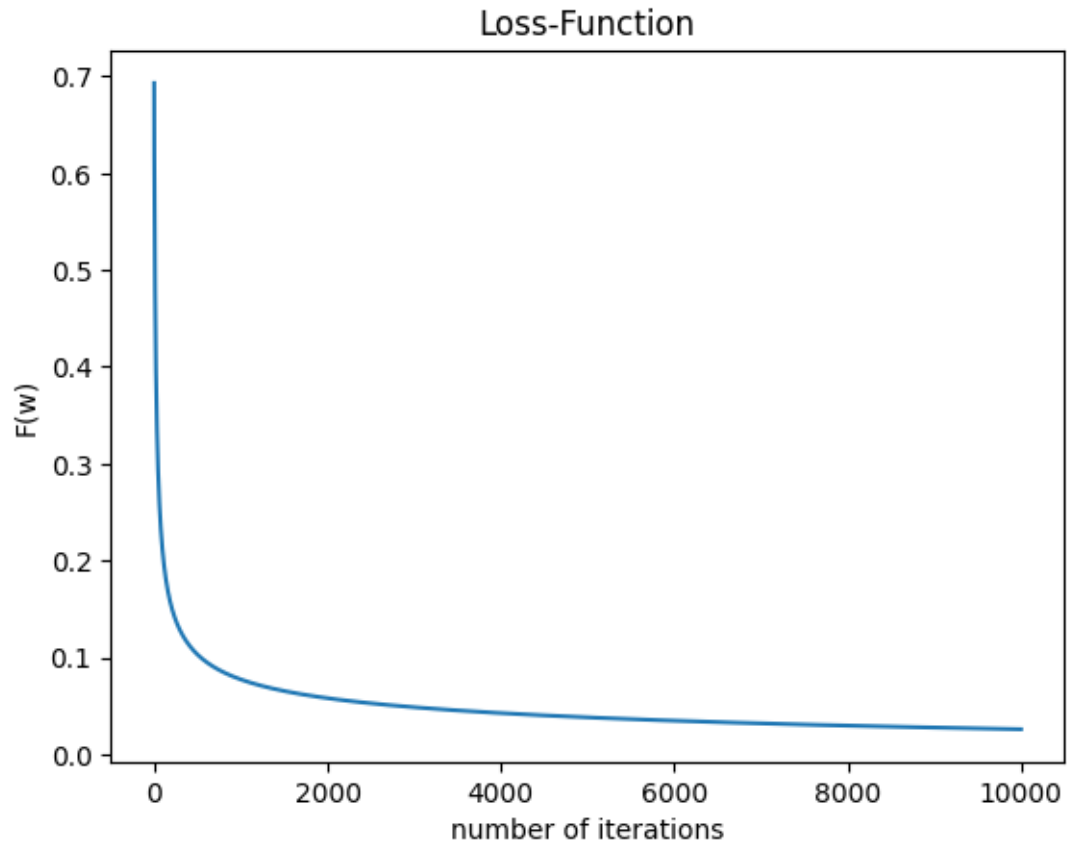
def armijo_condition(x:np.ndarray, lr:float, gamma,grad) -> bool:
    left = function(x + lr*grad,map)
    right = function(x,map) - gamma * lr * np.linalg.norm(grad)**2
    return left <= right

def backtracking(x: np.ndarray, iterations:int=iteration_nbr, gamma=0.5, beta=0.
↪8, lr=0.1) -> np.ndarray:
    x_values = [0]
    y_values = [function(x, map)]
    for i in range(int(iterations)):
        printf(weights, map, n)
        learning_rate = lr
        grad = optimized_gradient(x,map)
        function_value = function(x,map)
        while not armijo_condition(x,lr=learning_rate, gamma=gamma,grad=grad):
            learning_rate *= beta
        x -= lr * gamma * grad
        x_values.append(i)
        y_values.append(function(x,map))
        #print(f'x: {x}, f(x): {function(x,map)}')
        printf(weights, map, 10000)
    return x_values, y_values

start_value = np.array([0 for _ in range(784)], dtype=np.float64)
x_values, y_values = backtracking(start_value)

plt.plot(x_values, y_values)
plt.xlabel('number of iterations')
plt.ylabel('F(w)')
plt.title('Loss-Function')
#plt.xticks(range(1, iteration_nbr + 1, 500))
plt.savefig('F(w).png', format='png', dpi=300)
plt.show()
```

iteration : 10000, F(w) = 0.6931471805599322



5 Gradient descent with nesterov's acceleration

```
[36]: # 4 a
import math
lr = 1e-3
iteration_nbr = 10**4
beta = 1
xprev = np.array([0 for _ in range(784)], dtype=np.float64)

weights = np.array([0 for _ in range(784)], dtype=np.float64)

x_values = []
y_values = []

for n in range(iteration_nbr):
    printf(weights, map, n)
    x_values.append(n)
    y_values.append(function(weights, map))
    x_save = weights
```

```

    weights = weights - lr * optimized_gradient((weights + beta * (weights -
↪xprev)), map) + beta * (weights - xprev)
    xprev = x_save
printf(weights, map, 10000)

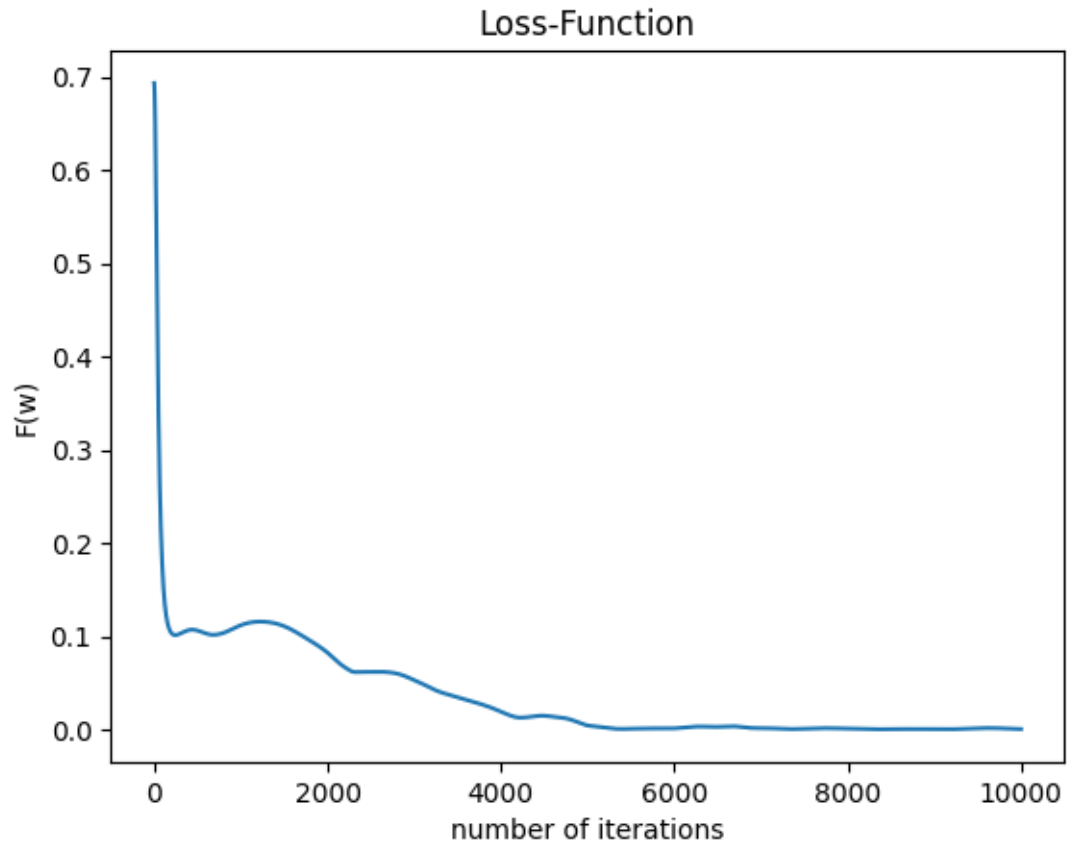
plt.plot(x_values, y_values)
plt.xlabel('number of iterations')
plt.ylabel('F(w)')
plt.title('Loss-Function')
plt.ticklabel_format(style='plain', axis='y') # Plain formatting for y-axis
#plt.xticks(range(1, iteration_nbr + 1, 500))
plt.savefig('F(w).png', format='png', dpi=300)
plt.show()

```

```

iteration : 0, F(w) = 0.6931471805599322
iteration : 1000, F(w) = 0.11172350057316358
iteration : 2000, F(w) = 0.08227191594481818
iteration : 3000, F(w) = 0.052877920857782625
iteration : 4000, F(w) = 0.018923473819430183
iteration : 5000, F(w) = 0.004191385416169663
iteration : 6000, F(w) = 0.0011126185653139691
iteration : 7000, F(w) = 0.0013404244476149909
iteration : 8000, F(w) = 0.0008098591758488051
iteration : 9000, F(w) = 0.00024442335332544314
iteration : 1000, F(w) = 0.00034703225744883584

```



6 Gradient descent with momentum

```
[37]: # 4 a
import math
lr = 1e-3
iteration_nbr = 10**4
beta = 1
xprev = np.array([0 for _ in range(784)], dtype=np.float64)

weights = np.array([0 for _ in range(784)], dtype=np.float64)

x_values = []
y_values = []

for n in range(iteration_nbr):
    printf(weights, map, n)
    x_values.append(n)
    y_values.append(function(weights, map))
    x_save = weights
```



```

    weights = weights - lr * optimized_gradient(weights, map) + beta * (weights_
↪- xprev)
    xprev = x_save
printf(weights, map, 10000)

plt.plot(x_values, y_values)
plt.xlabel('number of iterations')
plt.ylabel('F(w)')
plt.title('Loss-Function')
#plt.xticks(range(1, iteration_nbr + 1, 500))
plt.savefig('F(w).png', format='png', dpi=300)
plt.show()

```

```

iteration : 0, F(w) = 0.6931471805599322
iteration : 1000, F(w) = 0.11431282483016972
iteration : 2000, F(w) = 0.08673958820295308
iteration : 3000, F(w) = 0.05629395487577995
iteration : 4000, F(w) = 0.02191169896758379
iteration : 5000, F(w) = 0.006880022811634531
iteration : 6000, F(w) = 0.0010065529945642447
iteration : 7000, F(w) = 0.0013990424544088028
iteration : 8000, F(w) = 0.0011878032433038268
iteration : 9000, F(w) = 0.0005675642487008104
iteration : 1000, F(w) = 0.0008088792810672413

```

